



Parallel ML

Maciej Marchwiany, PhD



<https://www.linkedin.com/pulse/ai-powered-physics-how-ai-revolutionizing-scientific-discovery-cain-ddbwc>

Plan



Scikit learn



Data
parallelization



Model
parallelization



Horovod

Scikit learn



Parallel algorithms

1. Ensemble Methods

1. Random Forest - Uses parallelism for training multiple decision trees.

1. `RandomForestClassifier`

2. `RandomForestRegressor`

2. Gradient Boosting - Parallelizes only over data samples, not trees.

1. `GradientBoostingClassifier`

2. `GradientBoostingRegressor`

3. Histogram-Based Gradient Boosting - Uses parallel histogram-based training.

1. `HistGradientBoostingClassifier`

2. `HistGradientBoostingRegressor`

4. Extra Trees - Uses parallelism for training multiple decision trees

1. `ExtraTreesClassifier`

2. `ExtraTreesRegressor`

Parallel algorithms

1. Nearest Neighbors

1. K-Neighbors - Uses parallelism for distance computations.

1. `KNeighborsClassifier`

2. `KNeighborsRegressor`

Parallel algorithms

1. Linear Models

1. Ridge Regression with cross-validation - Uses threads parallelism.

1. RidgeCV

Parallel algorithms

1. Clustering

1. K-means - Uses threads parallelizm.

1. KMeans

2. MiniBatchKMeans

Parallel algorithms

1. Dimensionality Reduction

1. TruncatedSVD - parallel SVD decomposition.

1. TruncatedSVD

Parallel algorithms

1. Cross-Validation & Hyperparameter Tuning

1. Cross-Validation.

1. `cross_val_score`

2. Grid Search Cross-Validation.

1. `GridSearchCV`

3. Randomized Search Cross-Validation.

1. `RandomizedSearchCV`

Parallel algorithms

Algorithm	Parallel Support?	Notes
RandomForestClassifier/Regressor	✓ Full	Uses multiple trees
ExtraTreesClassifier/Regressor	✓ Full	Faster than Random Forest
HistGradientBoostingClassifier/Regressor	✓ Full	More efficient than GBM
KNeighborsClassifier/Regressor	✓ Full	Parallelizes distance computations
GradientBoostingClassifier/Regressor	✗ Limited	Only data parallelism
LinearSVC / LinearSVR	✗ Limited	Multi-threaded, no n_jobs
SVC / SVR	✗ No	No parallelism
DBSCAN / AgglomerativeClustering	✗ No	Single-threaded
KMeans	✓ Full	Uses OpenMP for multi-threading
PCA	✗ No	No parallelism
TruncatedSVD	✓ Full	n_jobs for parallel SVD
GridSearchCV / RandomizedSearchCV	✓ Full	Parallel hyperparameter tuning
cross_val_score	✓ Full	Parallel cross-validation

Parallel algorithm

```
import numpy as np.  
from sklearn.datasets import make_classification  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.model_selection import cross_val_score  
  
# Generate a synthetic dataset  
X, y = make_classification(n_samples=10000, n_features=20, random_state=42)  
  
# Initialize a Random Forest Classifier with parallelism  
rf = RandomForestClassifier(n_estimators=100, n_jobs=-1, random_state=42)  
  
# Perform 5-fold cross-validation in parallel  
scores = cross_val_score(rf, X, y, cv=5, n_jobs=-1)  
  
# Print results  
print(f"Cross-validation scores: {scores}")  
print(f"Mean accuracy: {np.mean(scores):.4f}")
```

Why parallel training?


FourCastNet: Large-compute scaling

Scaled to e.g. 3808 GPUs on Perlmutter with model parallel on 4-gpus


Kurth et al. 2022

[arXiv:2208.05419](https://arxiv.org/abs/2208.05419)

[PASC 2023 Plenary Video](#)



Train large models on ~1hr timescales compared to 40 hrs on 32 nodes or >~45days on a single GPU:




Model and weights made available to community at

<https://github.com/NVlabs/FourCastNet>



Data
parallelization



Data parallelization

Data parallelization is a technique where large datasets are **split across multiple processors or machines**, allowing computations (such as model training or inference) to happen in parallel. This helps speed up training and processing times, especially when dealing with **big data**.

How Data Parallelization Works

1. **Dataset Splitting** → The dataset is divided into smaller batches, each sent to a different worker (GPU, CPU, or cluster node).
2. **Model Duplication** → The **entire model** is copied onto each worker.
3. **Parallel Training** → Each worker **independently computes gradients** on its assigned mini-batch.
4. **Gradient Aggregation** → After each forward & backward pass, gradients from all workers are **averaged and synchronized**.
5. **Model Update** → The **updated parameters** are sent back to all workers, ensuring all models stay synchronized.

This technique is especially useful in **Scikit-Learn**, **Deep Learning (PyTorch, TensorFlow)**, and **distributed computing frameworks like Dask or Spark**.

Synchronization

- ▶ Synchronous updates
 - stable convergence
 - can be decentralized (allreduce)
 - computation may be blocked by communication
- ▶ Asynchronous updates
 - no waiting for gradients
 - stale gradients affect convergence
 - parameter server can be a bottleneck
- ▶ Delayed-synchronous updates
 - Lagged gradients allow better comms overlap
 - stale gradients affect convergence

When to Use Data Parallelization?

- ▶ When working with **large datasets**
- ▶ When training **complex models** that require **distributed computing**
- ▶ When using **multiple CPUs or GPUs** to speed up computation

PyTorch

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset

# Sample data
X = torch.randn(10000, 10)
y = torch.randint(0, 2, (10000,))

# Create a simple neural network
class SimpleNN(nn.Module):
    def __init__(self):
        super(SimpleNN, self).__init__()
        self.fc = nn.Linear(10, 2)

    def forward(self, x):
        return self.fc(x)

# Use multiple GPUs for training
model = SimpleNN()
model = nn.DataParallel(model)  # Enables multi-GPU parallelism

# Train with DataLoader (parallelized batching)
train_loader = DataLoader(
    TensorDataset(X, y), batch_size=32,
    shuffle=True, num_workers=4)

# Training loop
optimizer = optim.Adam(model.parameters())
loss_fn = nn.CrossEntropyLoss()
```

PyTorch

```
class SimpleNN(nn.Module):
    def __init__(self):
        super(SimpleNN, self).__init__()
        self.fc = nn.Linear(10, 2)

    def forward(self, x):
        return self.fc(x)

# Use multiple GPUs for training
model = SimpleNN()
model = nn.DataParallel(model)  # Enables multi-GPU parallelism

# Train with DataLoader (parallelized batching)
train_loader = DataLoader(
    TensorDataset(X, y), batch_size=32,
    shuffle=True, num_workers=4)

# Training loop
optimizer = optim.Adam(model.parameters())
loss_fn = nn.CrossEntropyLoss()

for epoch in range(5):
    for batch_X, batch_y in train_loader:
        optimizer.zero_grad()
        outputs = model(batch_X)
        loss = loss_fn(outputs, batch_y)
        loss.backward()
        optimizer.step()
```

Limitations

For large-scale deep learning, traditional Data Parallelism (like `torch.nn.DataParallel`) has limitations, such as:

- ▶ Memory inefficiency (each GPU stores the full model).
- ▶ Slow communication overhead (synchronizing gradients across GPUs).

Advanced techniques solve these issues by optimizing memory usage, reducing communication overhead, and improving scalability.

FSDP

Fully Sharded Data Parallelism

FSDP shards both model parameters and gradients across GPUs, unlike traditional Data Parallelism, which keeps full copies.

Key Benefits

- ▶ **Lower Memory Usage** - Each GPU stores only a part of the model.
- ▶ **Scalability** - Trains LLMs like GPT-3, LLaMA on multiple GPUs.
- ▶ **Gradient Communication Efficiency** - Reduces sync time.

FSDP

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.distributed.fsdp import FullyShardedDataParallel
as FSDP
from torch.distributed.fsdp.fully_sharded_data_parallel
import CPUOffload

# Initialize distributed process group
torch.distributed.init_process_group(backend="nccl")

# Define a simple model
class Transformer(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(1024, 4096)
        self.fc2 = nn.Linear(4096, 1024)

    def forward(self, x):
        return self.fc2(torch.relu(self.fc1(x)))

# Move model to FSDP (fully sharded)
device = torch.device("cuda")
model = Transformer().to(device)
```

FSDP

```
        return self.lc2(torch.relu(self.lc1(x)))

# Move model to FSDP (fully sharded)
device = torch.device("cuda")
model = Transformer().to(device)
model = FSDP(model,
cpu_offload=CPUOffload(offload_params=True)) # Moves params
to CPU to save memory

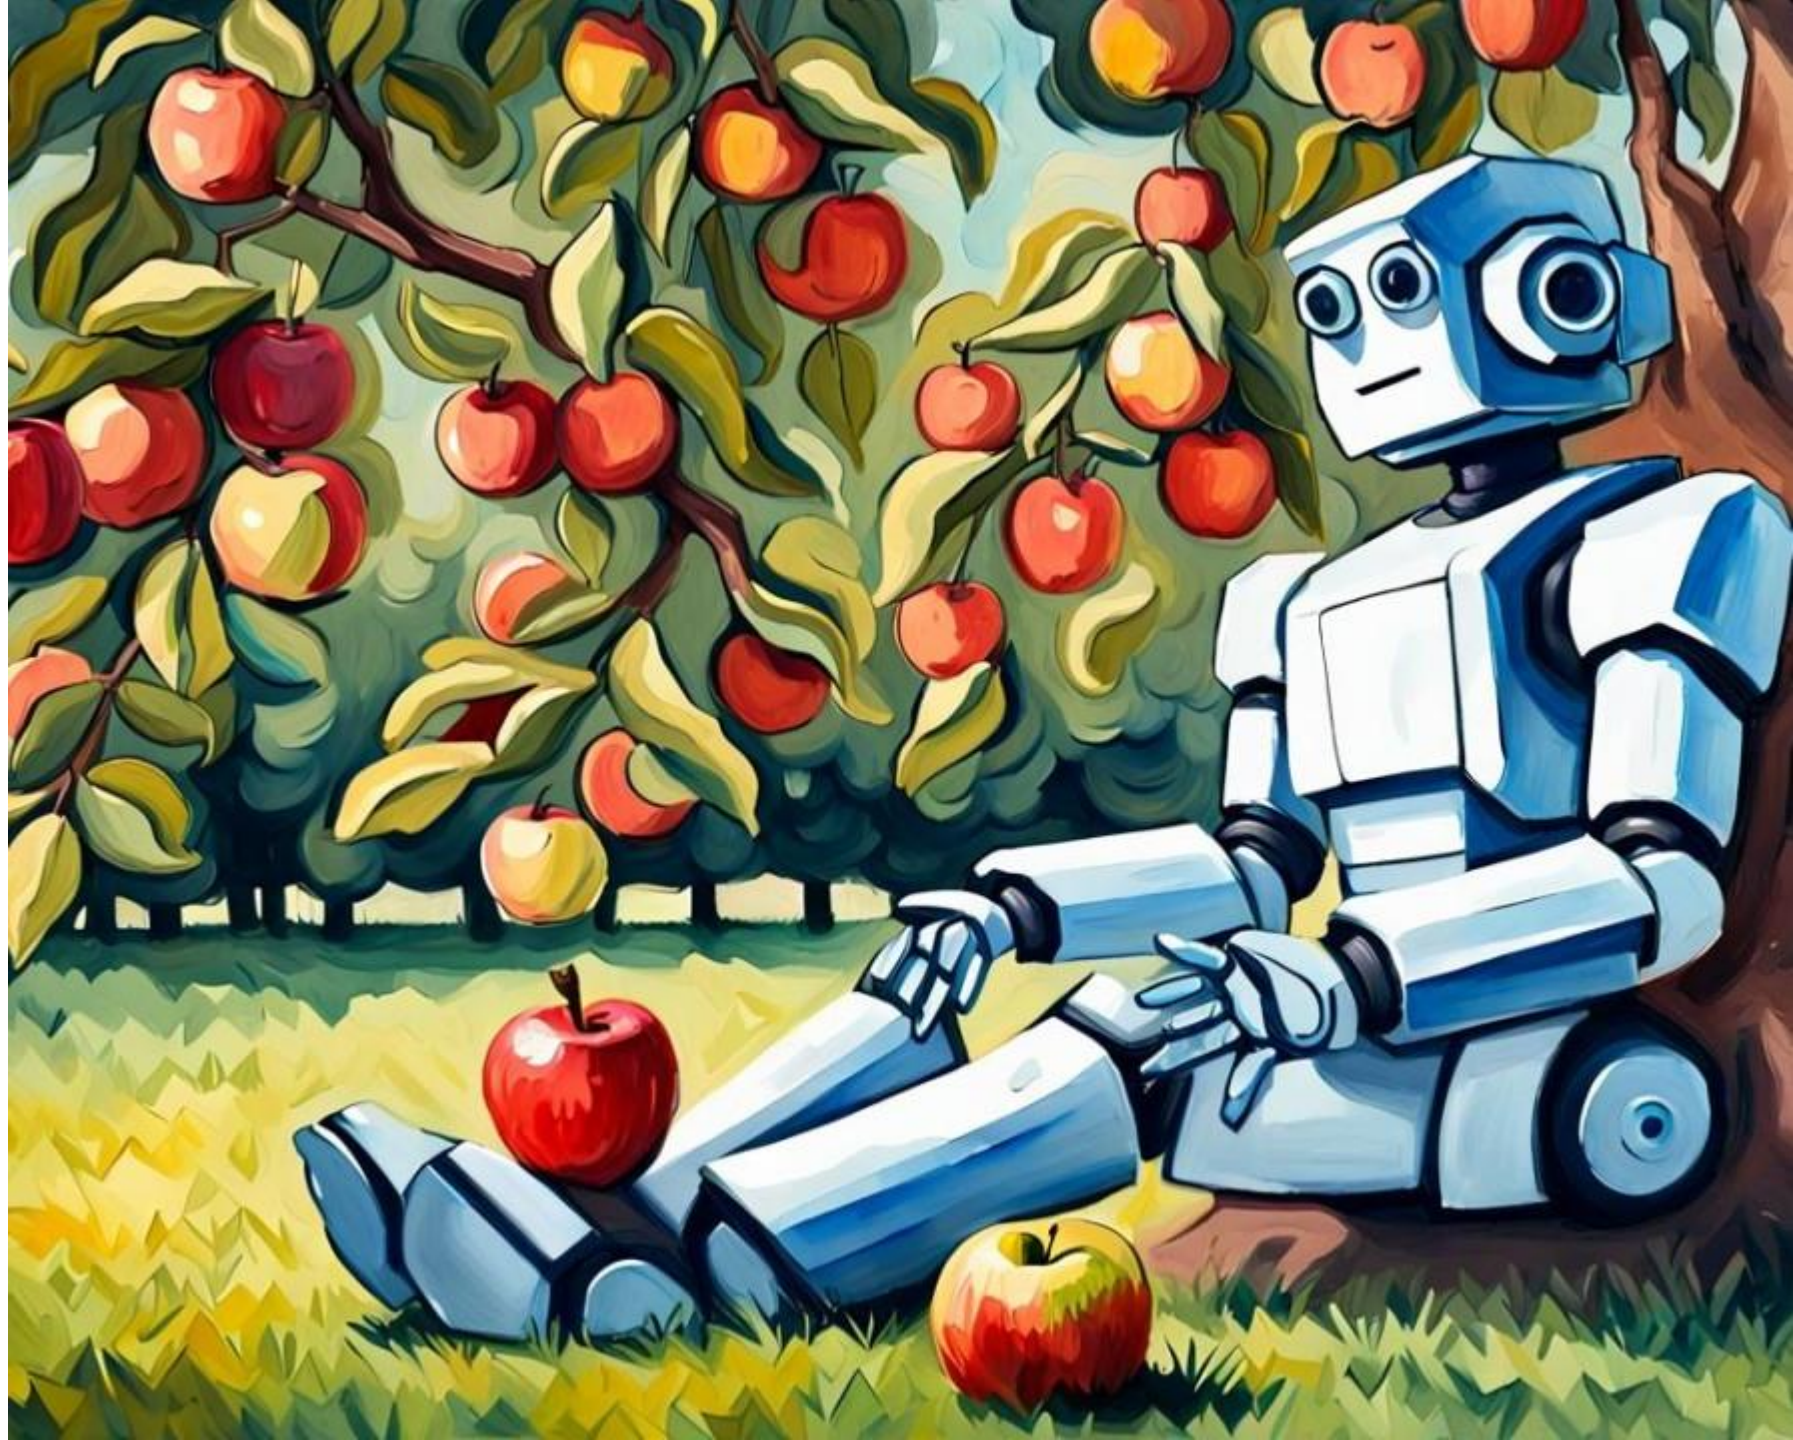
# Define optimizer & loss
optimizer = optim.AdamW(model.parameters(), lr=3e-4)
criterion = nn.MSELoss()

# Dummy data for training
x = torch.randn(32, 1024).to(device)
target = torch.randn(32, 1024).to(device)

# Forward & backward pass
output = model(x)
loss = criterion(output, target)
loss.backward()
optimizer.step()

print(f"Training done with FSDP! Loss: {loss.item()}")
```

Model
parallization



Model parallelization

Model parallelization is a technique where **different parts of a model are distributed across multiple processors, GPUs, or machines** to handle large-scale computations that a single device cannot manage efficiently.

How Model Parallelization Works

1. **Splitting the Model:** Different layers or parts of the model are assigned to different processors/GPUs.
2. **Parallel Computation:** Each processor computes only the assigned portion of the model.
3. **Communication & Synchronization:** The outputs of one stage (or layer) are transferred to the next stage, ensuring correct computations.

Model parallelization is mostly used in **deep learning**, where large models like GPT, BERT, or ResNets exceed a single device's memory capacity.

Types of Model Parallelism

- ▶ **Tensor Parallelism (TP) - Splitting Matrices Across GPUs**
 - Splits individual tensors (weights, activations, gradients) across multiple GPUs.
 - Common in Transformer models (GPT-3, LLaMA, ViTs), where matrix multiplications are expensive.
 - A single layer is spread across multiple GPUs, so each device computes only part of the operation.
- ▶ **Pipeline Parallelism - Splitting Model Layers Across GPUs**
 - Divides different layers of the model across GPUs.
 - Each GPU performs a stage of forward and backward computation.
- ▶ **Expert Parallelism (Mixture of Experts, MoE)**
 - Divides the model into multiple expert networks, where each GPU only computes a subset of the experts.
 - Only a few experts are active per forward pass → Huge memory savings!
 - Used in GPT-4, GLaM, Switch Transformer, and DeepSpeed MoE.

Tensor Parallelism

```
from megatron import initialize_megatron, get_args
from megatron.model import GPTModel

# Initialize Megatron
initialize_megatron(extra_args_provider=None)

# Set model parallelism (split tensors across GPUs)
args = get_args()
args.tensor_model_parallel_size = 2 # Use 2 GPUs for
tensor parallelism

# Define GPT model with tensor parallelism
model = GPTModel(num_layers=24, hidden_size=1024,
num_attention_heads=16)

print("Model initialized with Tensor Parallelism!")
```

Pipeline Parallelism

```
import torch
import torch.nn as nn

# Define a model with different layers assigned to different GPUs
class ModelParallelNN(nn.Module):
    def __init__(self):
        super(ModelParallelNN, self).__init__()
        # First layer on GPU 0
        self.layer1 = nn.Linear(1024, 512).to('cuda:0')
        # Second layer on GPU 1
        self.layer2 = nn.Linear(512, 256).to('cuda:1')
        # Third layer on GPU 1
        self.layer3 = nn.Linear(256, 10).to('cuda:1')

    def forward(self, x):
        x = x.to('cuda:0')    # Send input to GPU 0
        x = self.layer1(x)
        x = x.to('cuda:1')    # Move output to GPU 1
        x = self.layer2(x)
        x = self.layer3(x)
        return x

# Create model
model = ModelParallelNN()

# Generate sample input
x = torch.randn(64, 1024).to('cuda:0')    # Batch of 64 samples

# Forward pass
output = model(x)
print(output.shape)
```

Pipeline Parallelism

```
import torch
import torch.nn as nn
from torch.distributed.pipeline.sync import Pipe

# Define a simple deep model
class DeepModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.block1 = nn.Sequential(nn.Linear(1024, 2048), nn.ReLU())
        self.block2 = nn.Sequential(nn.Linear(2048, 1024), nn.ReLU())

    def forward(self, x):
        x = self.block1(x)
        return self.block2(x)

# Split model into pipeline stages across 2 GPUs
model = DeepModel()
model = Pipe(model, chunks=2, balance=[1, 1], devices=[0, 1])

print("Model initialized with Pipeline Parallelism!")
```

Expert Parallelism

```
import torch
import torch.nn as nn

class MixtureOfExperts(nn.Module):
    def __init__(self, num_experts=4):
        super().__init__()
        self.experts = nn.ModuleList(
            [nn.Linear(1024, 1024)
             for _ in range(num_experts)])
        self.gate = nn.Linear(1024, num_experts) # Gating function

    def forward(self, x):
        gate_scores = torch.softmax(self.gate(x), dim=-1)
        output = sum(gate_scores[:, i].unsqueeze(1)
                     * self.experts[i](x) for i in range(len(self.experts)))
        return output

model = MixtureOfExperts()
print("MoE Model Initialized!")
```


Hybrid Parallelism

For massive models like GPT-4, LLaMA-65B, PaLM-540B, we combine multiple parallelization techniques:

- ▶ **Data Parallelism** - Distribute data across GPUs.
- ▶ **Tensor Parallelism** - Split tensors across GPUs.
- ▶ **Pipeline Parallelism** - Distribute layers across GPUs.
- ▶ **Mixture of Experts (MoE)** - Activate only a subset of experts per forward pass.

Horovod





Bottleneck



Bottlenecks for ML speed

1. Computational Power and Hardware Limitations
 - **GPU and CPU Bottlenecks:** While GPUs are powerful for ML computations, they can be underutilized due to slow data input/output (I/O) operations. CPUs are often insufficient for handling pre-processing tasks efficiently, leading to bottlenecks in data preparation.
 - **Memory Bottlenecks:** The disparity between processing speeds and memory access speeds creates significant bottlenecks. As models become more complex, memory bandwidth becomes a limiting factor.
2. Data Ingestion and Storage
 - **I/O Bottlenecks:** The speed at which data can be loaded into memory is often slower than the computation speed of GPUs. This results in GPUs waiting for data, reducing overall efficiency.
 - **Storage Limitations:** Large datasets exceed DRAM capacity, causing I/O bottlenecks during training. Solutions like massively parallel storage systems are being developed to address this.
3. Software Optimization and Integration
 - **Software Frameworks:** Efficient software frameworks are crucial for maximizing AI inference performance. However, optimizing software for specific hardware configurations remains a challenge.
 - **Deployment Challenges:** A significant bottleneck is the transition from model development to production. This often involves disconnects between data scientists and IT teams.

Bottlenecks for ML speed

- **Storage Limitations:** Large datasets exceed DRAM capacity, causing I/O bottlenecks during training. Solutions like massively parallel storage systems are being developed to address this.
3. **Software Optimization and Integration**
 - **Software Frameworks:** Efficient software frameworks are crucial for maximizing AI inference performance. However, optimizing software for specific hardware configurations remains a challenge.
 - **Deployment Challenges:** A significant bottleneck is the transition from model development to production. This often involves disconnects between data scientists and IT teams.
 4. **Data Preprocessing**
 - **CPU Preprocessing:** Preprocessing data using CPUs can be a major bottleneck. Solutions like offloading preprocessing to GPUs are being explored.